

# Homework 3: Neural Networks

CS412

Released: February 20th

Due: March 6th, 11:30pm on Gradescope

## 1 Neural Network with Backpropagation

The primary portion of this assignment will be coding a neural network that finds a solution using gradient descent and backpropagation. While you may use any language for the following code, you should discuss with the instructor if your choices will be other than Python (I think R will be very difficult here).

I have provided some basic structure, as well as **an example that shows how to build the 2-layer 4-neuron network explicitly**. You will need to allow three additional pieces of functionality:

- a)  $\eta$ : Your function should perform gradient descent relative to a learning rate
- b) MaxIterations: Your neural network should halt after it goes through the given number of iterations
- c) numLayers,numNeurons: numLayers should be the number of hidden layers and numNeurons is the number of hidden neurons per layer. You may assume that numLayers and numNeurons will be at least one.

## 2 Neural Network Testing

In this section, run package neural networks **on the digits data and on the new dataset**. **HERE** is the python library for the neural network classifier. For all models tested, you should use `activation = "logistic"` and `solver = "sgd"`. All other inputs can be default. **For our exercise, all hidden layers will have the same number of neurons.**

For the digits data, and your two features from HW1, give the cross validation errors for the following. Define the number of hidden layers to be 1-5 and the number of nodes per layer to be 2-4. Give the cross validation error for each of these. **A table is sufficient here.** **Report the number of hidden layers and nodes per layer which provides the highest accuracy.**

Plot the decision region for the optimal neural network above in the 2D space **Label this figure 3.1**

From the new dataset, define the number of hidden layers to be from 1-20 and the number of nodes per layer to be 2-20. For each of these 180 models, report the cross validation error and the runtime in milliseconds.

### Short answer

- a) How does runtime scale with the number of layers and the number of nodes per layer? Do each have a similar effect?
- b) What is the number of layers and nodes per layer that gives an optimum result?
- c) For your optimum model, try increasing and decreasing the learning rate from the default (0.001). Discuss the tradeoff here between runtime and accuracy?
- d) Is the neural network finding the same solution everytime? Why or why not? Does this have an impact on the expected fit?
- e) **Graduate Student Question:** Experiment with early stopping on neural networks that have a large number of internal nodes. Does this cause the model to under or over fit the data? Why?

## Extra Credit

Perform kPCA with an rbf kernel of gamma of ten on the new dataset. Use this reduced feature space to fit a neural network with 10 layers of 10 nodes each. Experiment with learning rate, max iteration and regularization to analyze their affect on fitting. Discuss your findings. Plot the decision region in the 2D space. You do not need to standardize your features, but you may need to increase your drawing area.

## Making your report

When you submit, there will be two submissions on gradescope. One for your pdf report and another for your zip file containing all your code. If you use a language other than Python, include a list of any packages you downloaded in your report.